

Job Title: AI Engineer
Location: San Francisco

Join our AI team to develop machine learning models and deploy them to production. You will work on cutting-edge NLP and computer vision projects.

Requirements:

- 3+ years in machine learning
- Python proficiency
- Experience with PyTorch or TensorFlow
- NLP and transformer models
- MLOps and model deployment
- Docker and Kubernetes

Nice to have:

- Research experience
- ML frameworks like scikit-learn
- Cloud deployment experience